

# LangChain

## Chat with your data

### Overview



LangChain

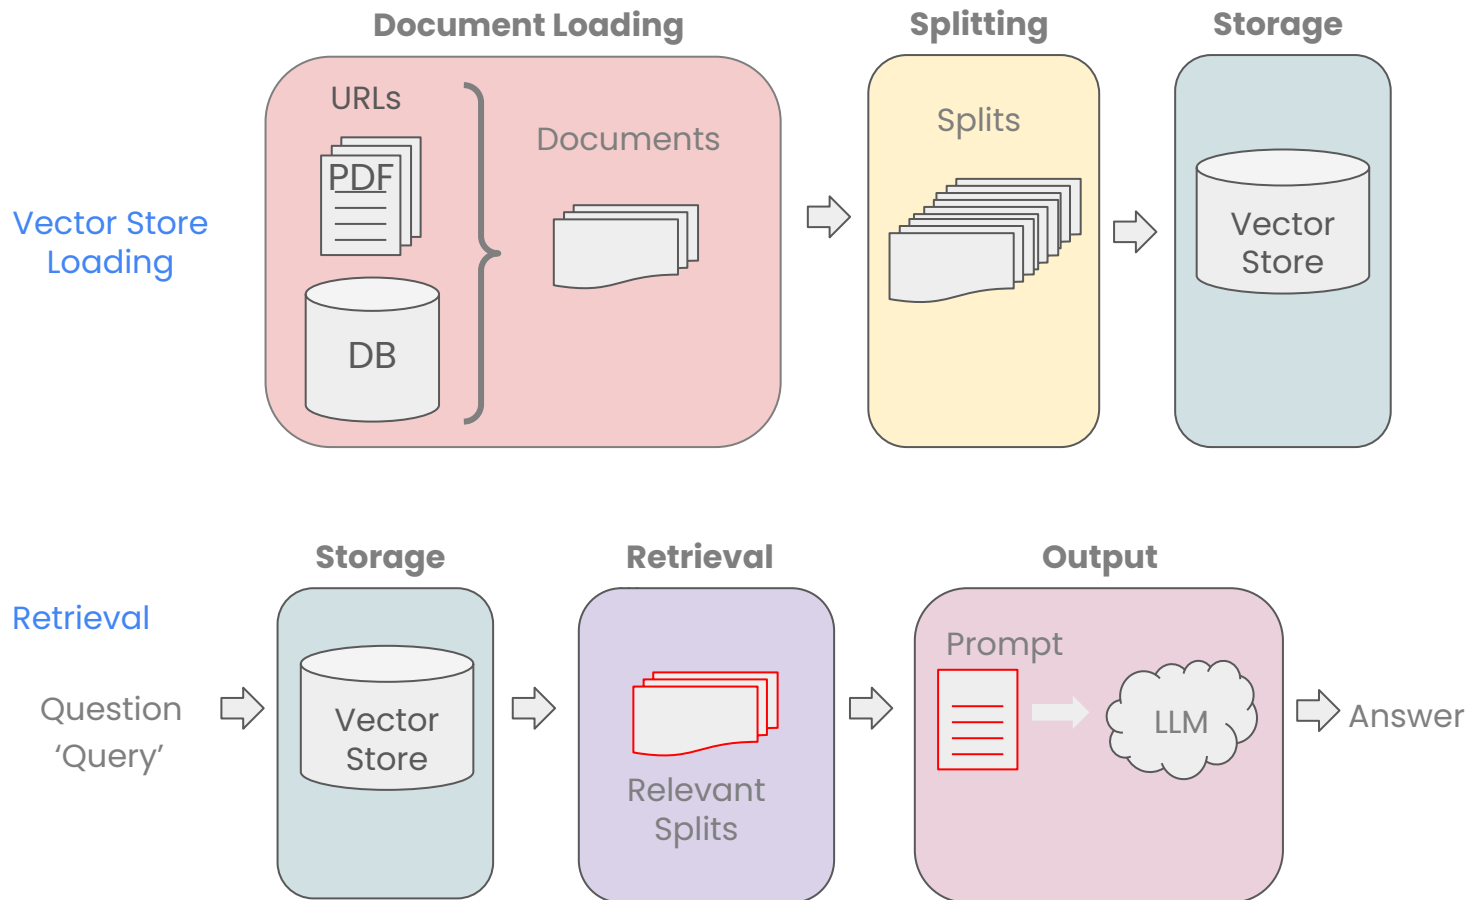


DeepLearning.AI

# Retrieval Augmented Generation

Retrieval Augmented Generation (RAG) is a very popular paradigm.

- Retrieve relevant documents and load into “working memory” / context window.



# LangChain

## Chat with your data

### Document Loading



LangChain



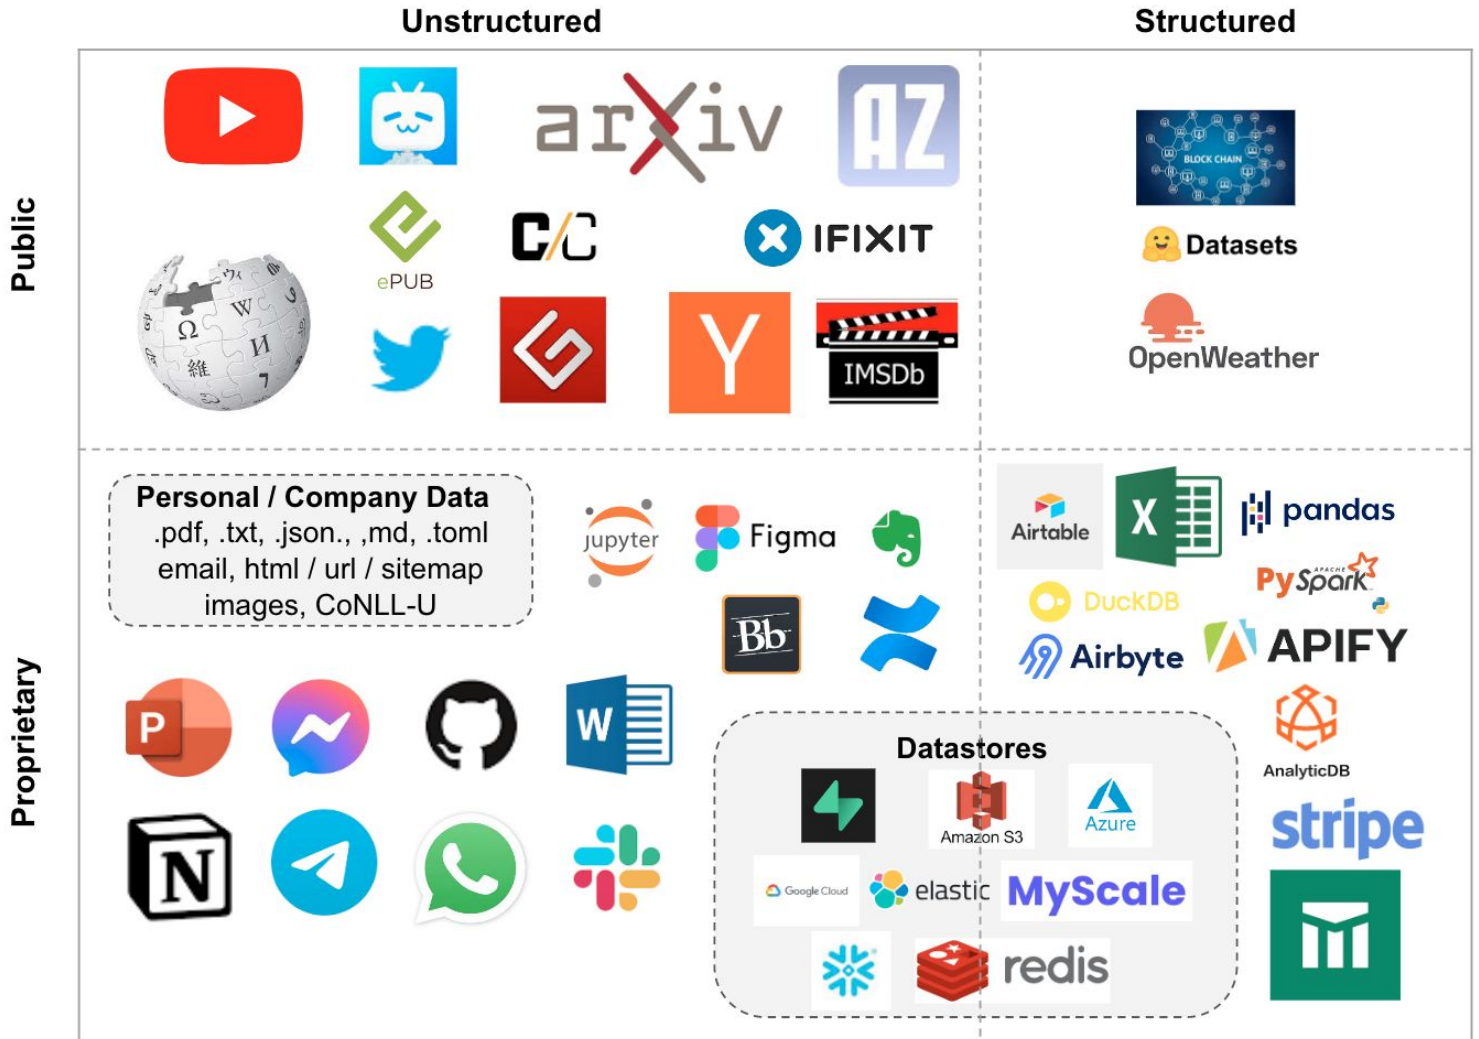
DeepLearning.AI

# Loaders

- Loaders deal with the specifics of accessing and converting data
  - Accessing
    - Web Sites
    - Data Bases
    - YouTube
    - arXiv
    - ...
  - Data Types
    - PDF
    - HTML
    - JSON
    - Word, PowerPoint...
- Returns a list of `Document` objects:

```
[
Document(page_content='MachineLearning-Lecture01 \nInstructor (Andrew Ng): Okay.
Good morning. Welcome to CS229....';
metadata={'source': 'docs/cs229_lectures/MachineLearning-Lecture01.pdf', 'page': 0})
...
Document(page_content='[End of Audio] \nDuration: 69 minutes ',
metadata={'source': 'docs/cs229_lectures/MachineLearning-Lecture01.pdf', 'page': 21})
]
```

# Document Loaders



# LangChain

## Chat with your data

### Document Splitting



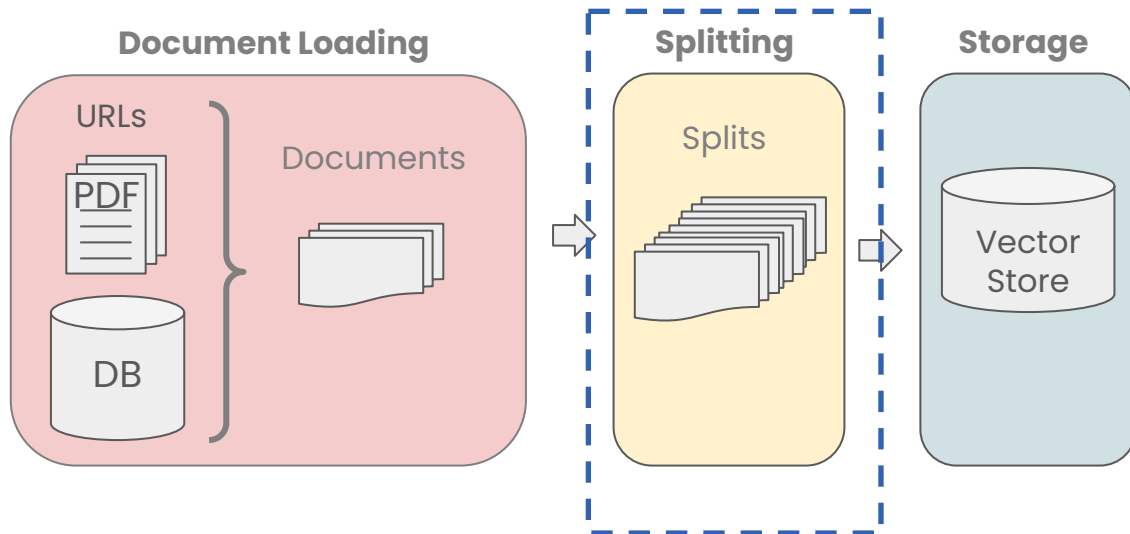
LangChain



DeepLearning.AI

# Document Splitting

- Splitting Documents into smaller chunks
  - Retaining meaningful relationships!



...  
 on this model. The Toyota Camry has a head-snapping  
 80 HP and an eight-speed automatic transmission that will  
 ...

**Chunk 1:** on this model. The Toyota Camry has a head-snapping

**Chunk 2:** 80 HP and an eight-speed automatic transmission that will

Question: What are the specifications on the Camry?

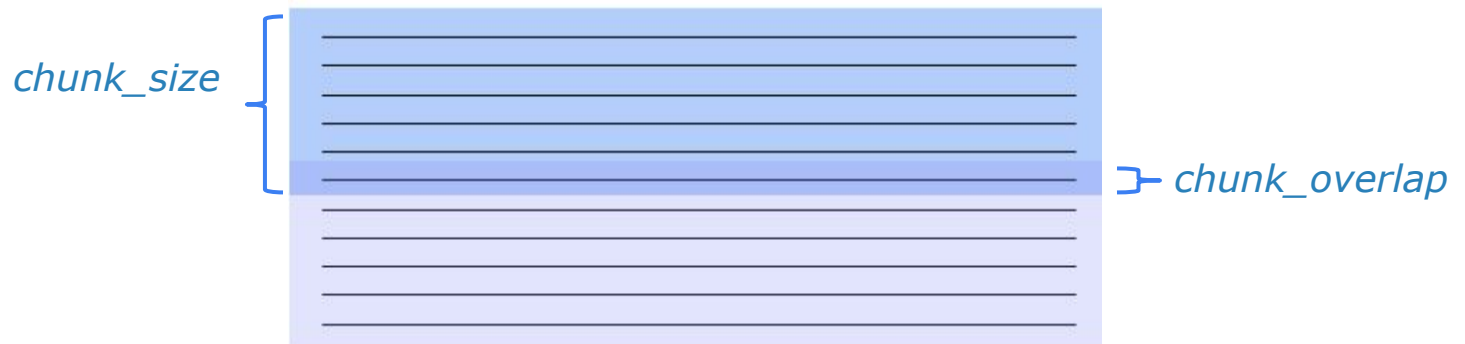
# Example Splitter

```
langchain.text_splitter.CharacterTextSplitter(  
    separator: str = "\n\n"  
    chunk_size=4000,  
    chunk_overlap=200,  
    length_function=<builtin function len>,  
)
```

Methods:

**create\_documents()** - Create documents from a list of texts.

**split\_documents()** - Split documents.





# Types of splitters

langchain.text\_splitter.

- **CharacterTextSplitter()**- Implementation of splitting text that looks at characters.
- **MarkdownHeaderTextSplitter()** - Implementation of splitting markdown files based on specified headers.
- **TokenTextSplitter()** - Implementation of splitting text that looks at tokens.
- **SentenceTransformersTokenTextSplitter()** - Implementation of splitting text that looks at tokens.
- ***RecursiveCharacterTextSplitter()*** - Implementation of splitting text that looks at characters. Recursively tries to split by different characters to find one that works.
- **Language()** – for CPP, Python, Ruby, Markdown etc
- **NLTKTextSplitter()** - Implementation of splitting text that looks at sentences using NLTK (Natural Language Tool Kit)
- **SpacyTextSplitter()** - Implementation of splitting text that looks at sentences using Spacy

# LangChain

## Chat with your data

Vector Stores and Embeddings

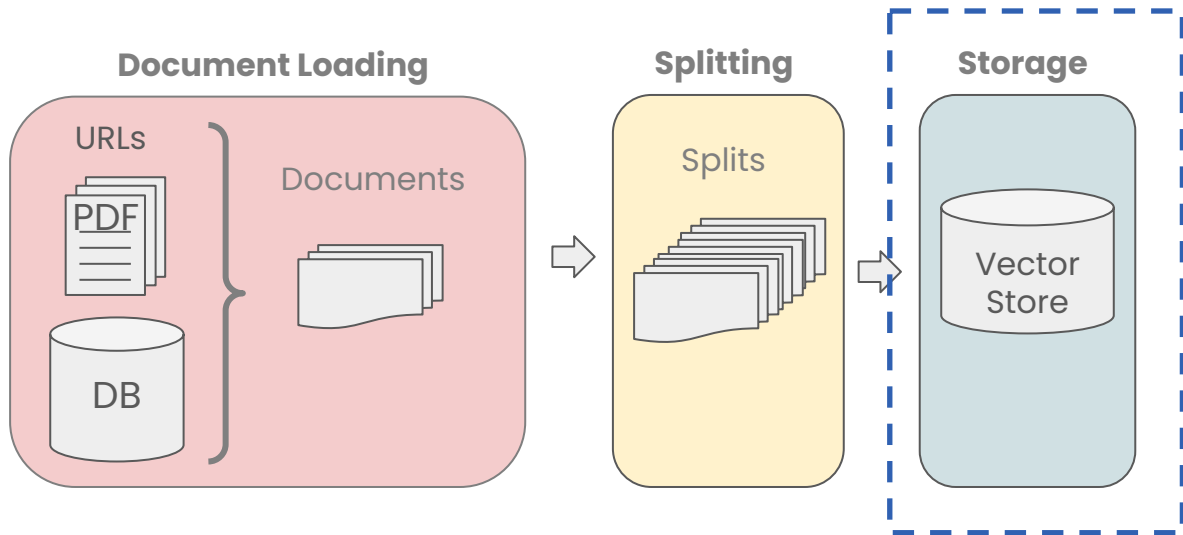


LangChain




DeepLearning.AI

# Vector Stores

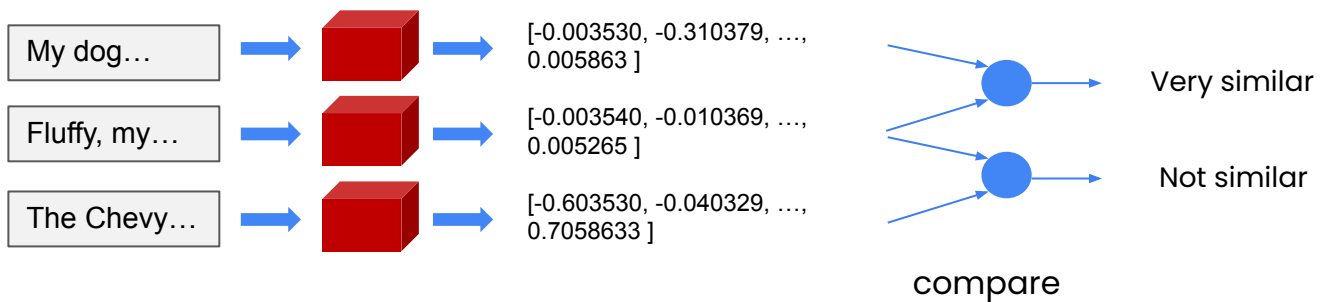


# Embeddings



- Embedding vector captures content/meaning
- Text with similar content will have similar vectors


- 1) My dog Rover likes to chase squirrels.
- 2) Fluffy, my cat, refuses to eat from a can.
- 3) The Chevy Bolt accelerates to 60 mph in 6.7 seconds.




# Vector Store

create

Vector Store

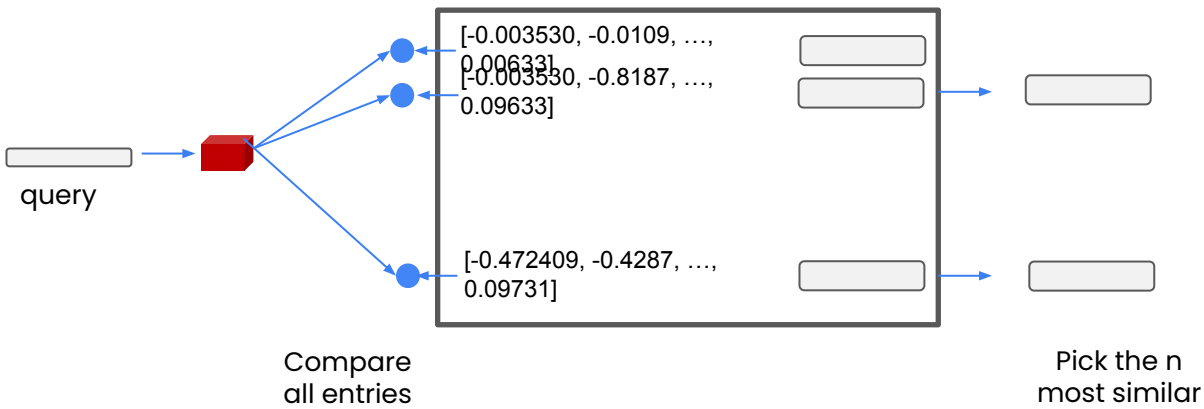


index




# Vector Store/Database

index



Process with llm



The returned values can now fit in the LLM context

# LangChain

## Chat with your data

### Retrieval

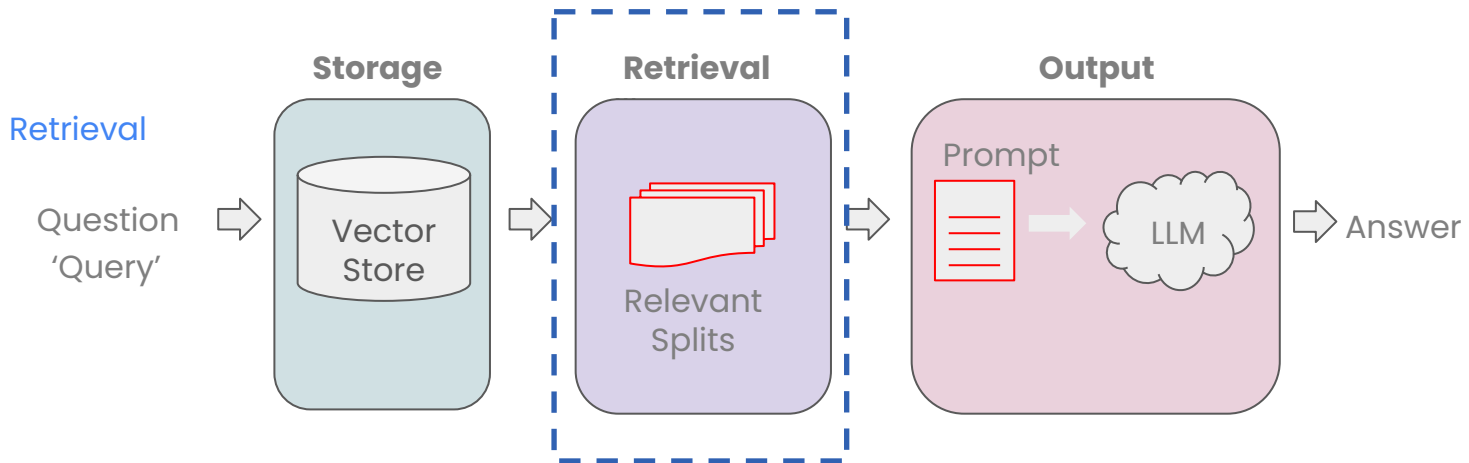


LangChain



DeepLearning.AI

# Retrieval



- Accessing/indexing the data in the vector store
  - Basic semantic similarity
  - Maximum marginal relevance
  - Including Metadata
- LLM Aided Retrieval



# Maximum marginal relevance(MMR)

- You may not always want to choose the most similar responses



Tell me about all-white mushrooms with large fruiting bodies


The Amanita phalloides has a large and imposing epigeous (aboveground) fruiting body (basidiocarp).

A mushroom with a large fruiting body is the Amanita phalloides. Some varieties are all-white.

AA. phalloides, a.k.a Death Cap, is one of the most poisonous of all known mushrooms.



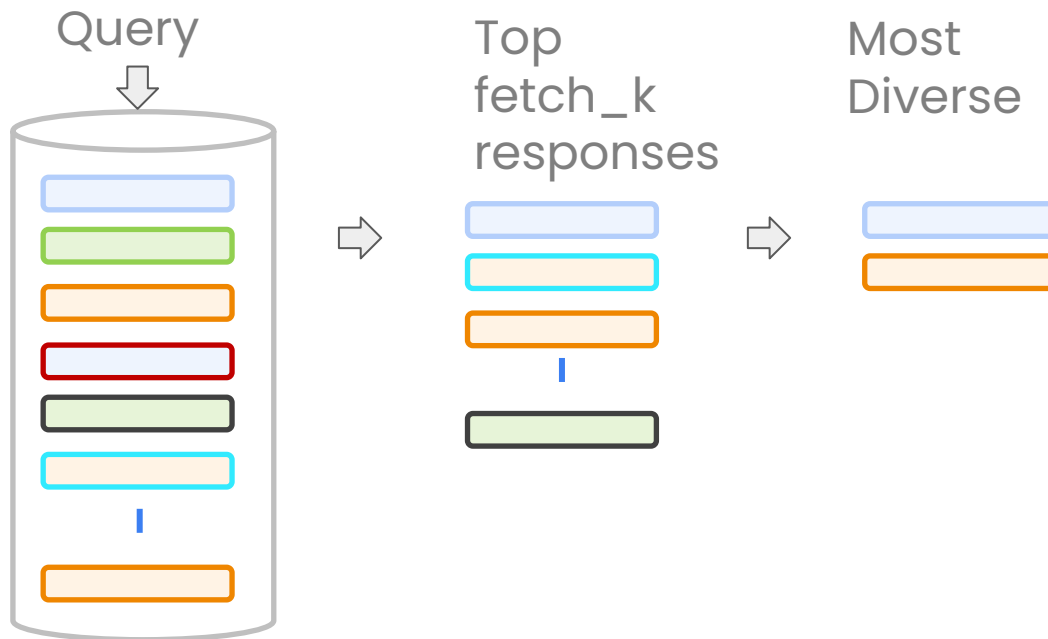
Most Similar



MMR

# MMR algorithm

- Query the Vector Store
- Choose the `fetch\_k` most similar responses
- Within those responses choose the `k` most diverse



# LLM Aided Retrieval

- There are several situations where the **Query** applied to the DB is more than just the **Question** asked.
- One is SelfQuery

## Self-query

Information + Question



LLM



Query



Post processing



Store



Relevant  
splits



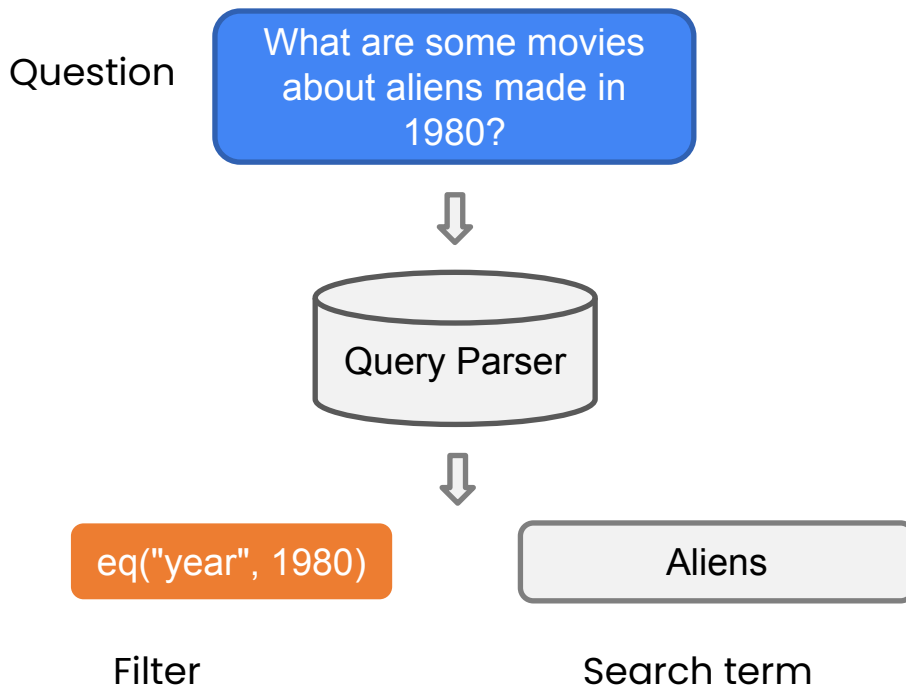
Information: Query format

Query: Question  
Filter: eq["section", "testing"]

Query parser


# LLM Aided Retrieval

- There are several situations where the **Query** applied to the DB is more than just the **Question** asked.
- One is SelfQuery, where we use an LLM to convert the user question into a query



# Compression

- Increase the number of results you can put in the context by shrinking the responses to only the relevant information.



# LangChain

## Chat with your data

### Question Answering

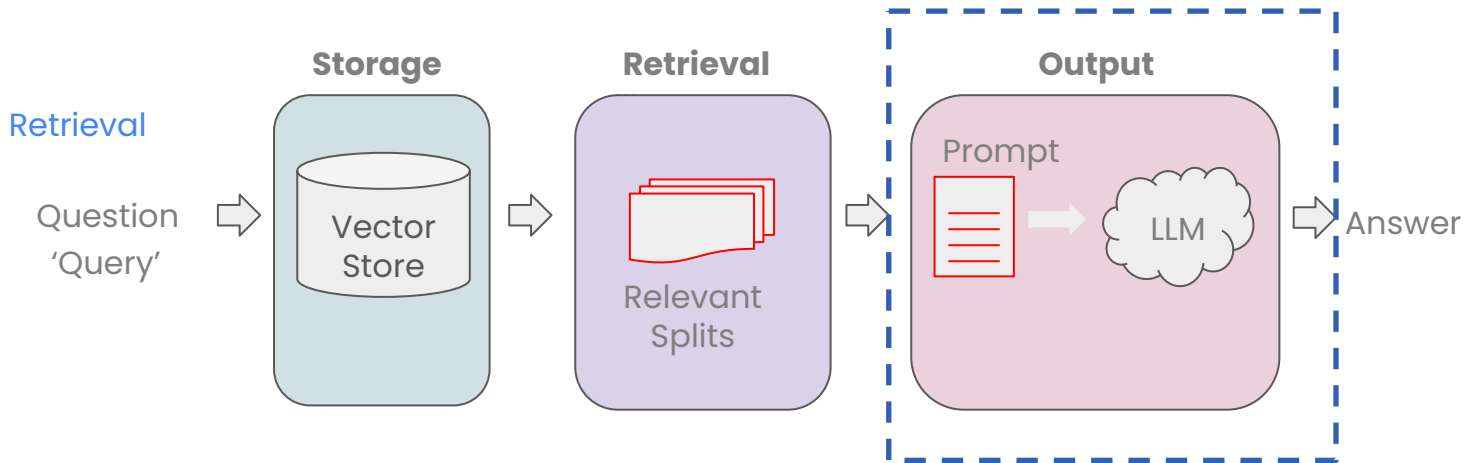


LangChain



DeepLearning.AI

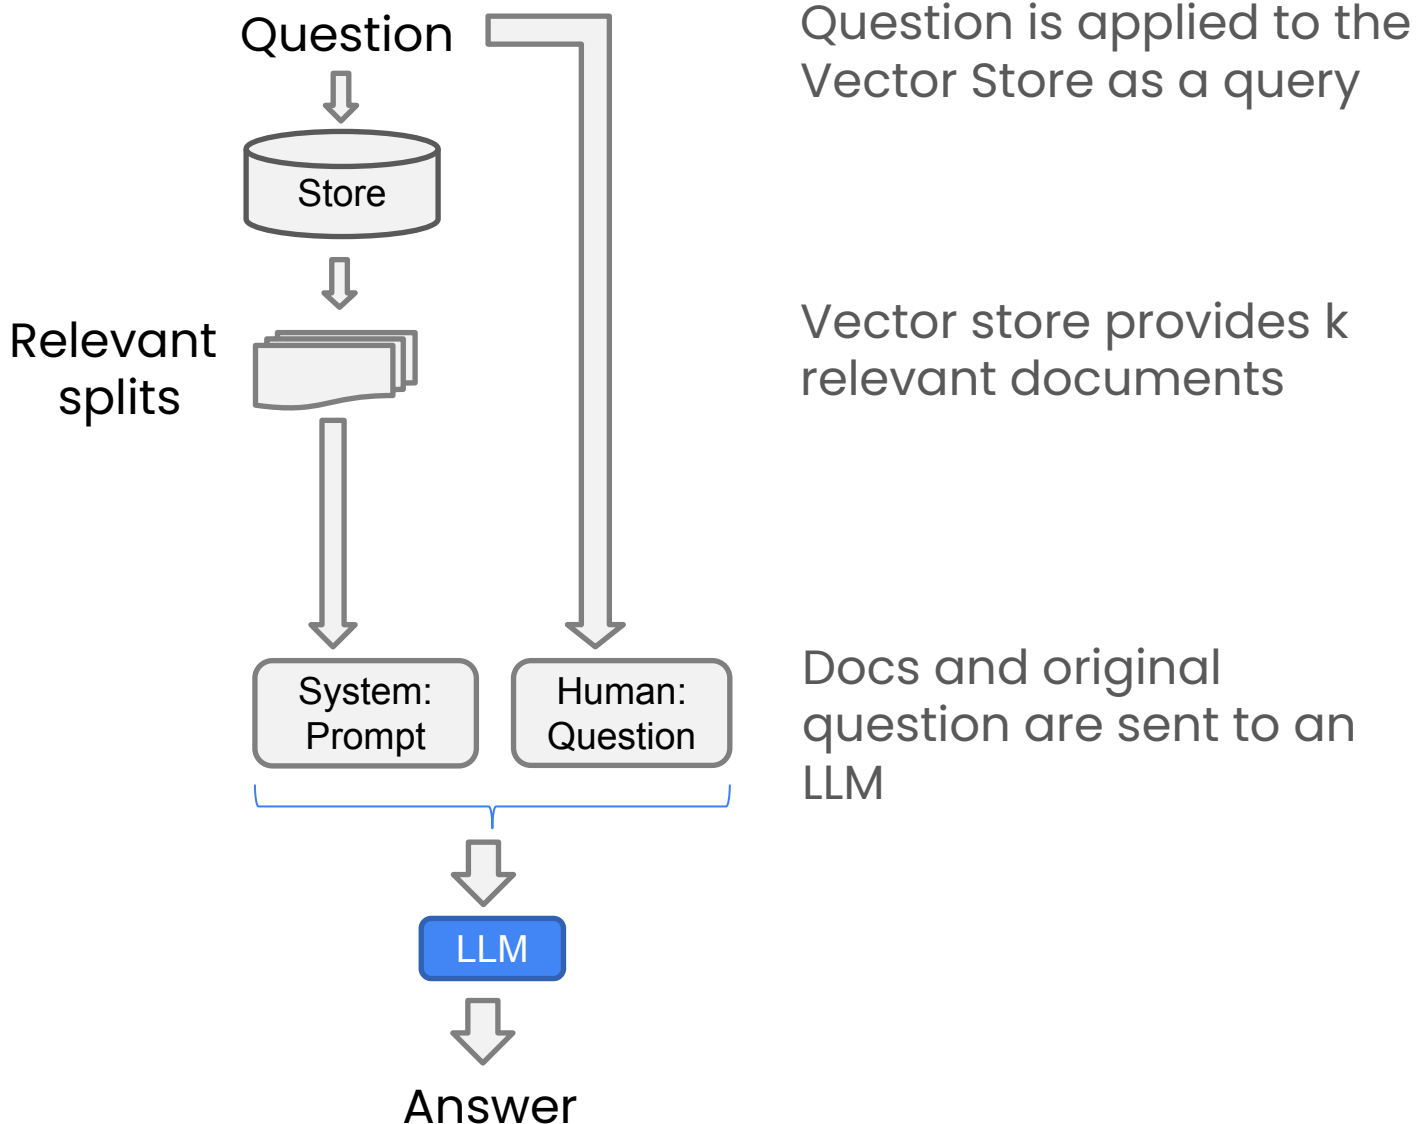
# Question Answering



- Multiple relevant documents have been retrieved from the vector store
- Potentially compress the relevant splits to fit into the LLM context
- Send the information along with our question to an LLM to select and format an answer

# RetrievalQA chain

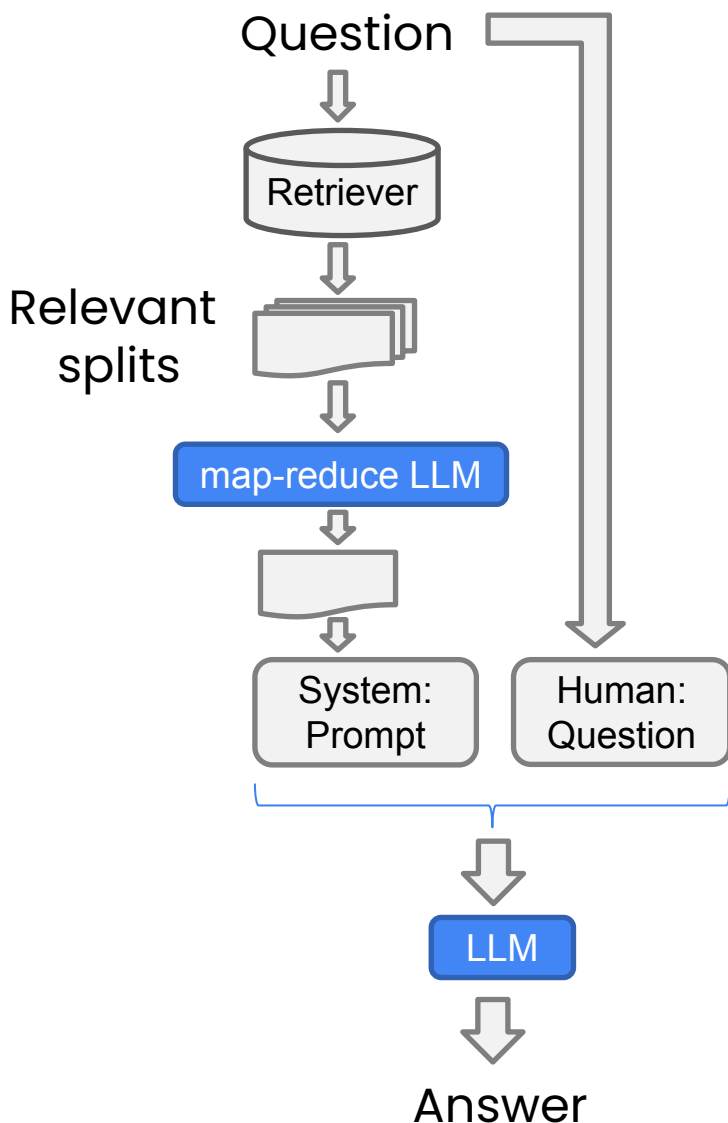
`RetrievalQA.from_chain_type(, chain_type="stuff", ...)`





# Retrieval Chain with LLM selection


`RetrievalQA.from_chain_type(, chain_type="map_reduce",...)`




You may have too many docs to fit into an LLM context. The solution is to use an LLM to select the 'most relevant' information

## 3 additional methods


### 1. Map\_reduce



### 2. Refine

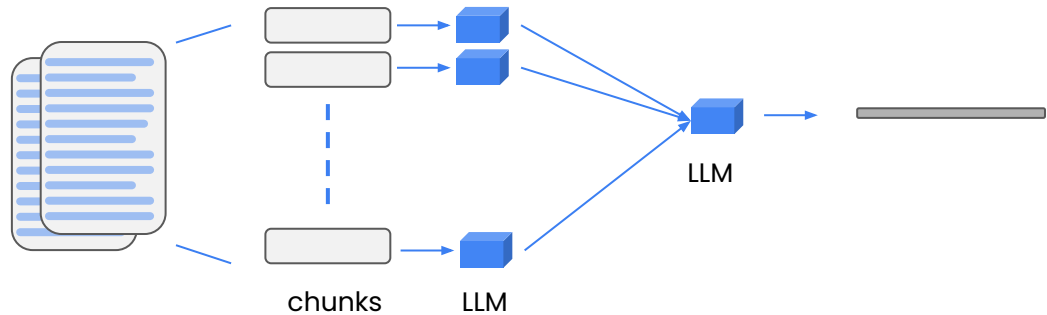


### 3. Map\_rerank




# 3 additional methods

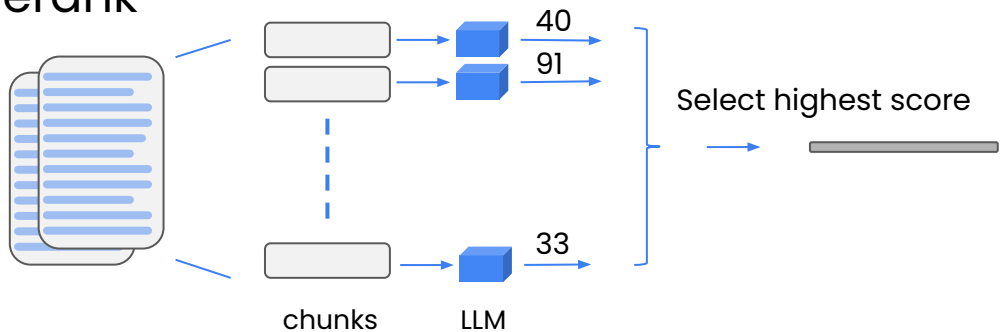
## 1. Map\_reduce



## 2. Refine



## 3. Map\_rerank



# LangChain for LLM Application Development

## Agents




LangChain




DeepLearning.AI

# Agents



Agent refers to the idea of using large language models as reasoning engine to determine which actions to take and in what order.

# Agents

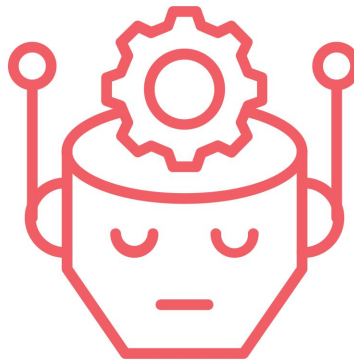


Agents use an LLM to determine which actions to take and in what order.

An action can be using a tool and observing its output and deciding what to return to the user.

```
from langchain.agents import  
initialize_agent, AgentType  
  
agent = initialize_agent(tools, llm,  
                        agent=AgentType.ZERO_SHOT_REACT_DESCRIPTION,  
                        verbose=True)
```

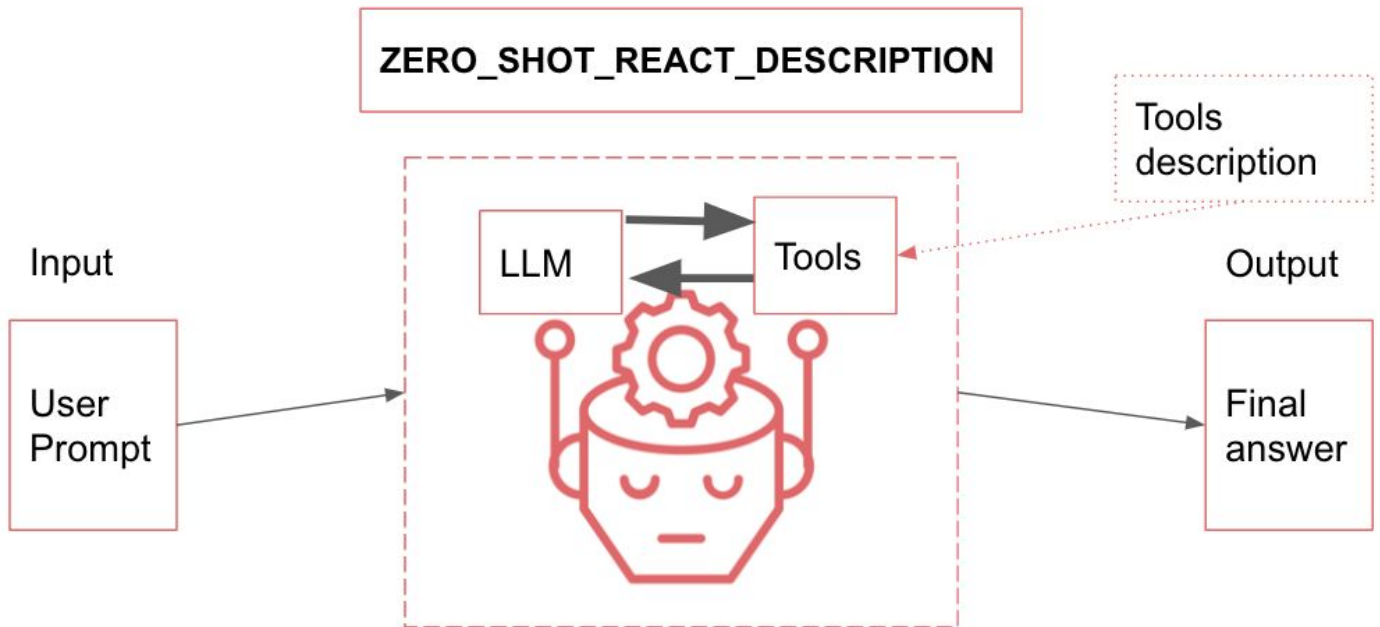
# Agents



To construct an agent, you need:

- **PromptTemplate:** this is responsible for taking the user input and previous steps and constructing a prompt to send to the language model
- **Language Model:** this takes the prompt constructed by the PromptTemplate and returns some output
- **Output Parser:** this takes the output of the Language Model and parses it into an AgentAction or AgentFinish object.

# Agents



# Agents

